

# Database Schema Manual

324827328, 213713522

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>The Core Entities</b>	<b>2</b>
2.1	Movies . . . . .	2
2.2	People . . . . .	2
<b>3</b>	<b>Connecting the Dots: Relationships</b>	<b>2</b>
3.1	The Cast & Crew . . . . .	2
3.2	Categorization (Genres, Keywords) . . . . .	3
<b>4</b>	<b>User Interaction</b>	<b>3</b>
4.1	Ratings . . . . .	3
<b>5</b>	<b>Binary Storage</b>	<b>3</b>
5.1	Storing Images in the DB . . . . .	3

# 1 Introduction

The database design was a critical component of the project, establishing a foundation for API performance and application integrity. We aimed for a normalized schema to minimize redundancy while maintaining the flexibility required for analytical queries.

## 2 The Core Entities

### 2.1 Movies

The `movies` table is the central entity of the database. We explicitly included `vote_average` and `vote_count` fields in this table.

This denormalization strategy was chosen to optimize read performance. Instead of calculating averages from the `ratings` table for every query, we maintain these aggregated values transactionally.

```
1 CREATE TABLE IF NOT EXISTS movies (
2     id INT PRIMARY KEY,
3     title VARCHAR(500) NOT NULL,
4     -- We keep the aggregated stats right here for speed
5     vote_average FLOAT,
6     vote_count INT,
7     -- ... other metadata
8     budget BIGINT,
9     revenue BIGINT,
10    status VARCHAR(50)
11 );
```

Listing 1: The heart of the system

### 2.2 People

We unified actors, directors, and writers into a single `people` table. This approach avoids data duplication for individuals who fulfill multiple roles (e.g., both acting and directing).

```
1 CREATE TABLE IF NOT EXISTS people (
2     id INT PRIMARY KEY,
3     name VARCHAR(255) NOT NULL,
4     gender INT,
5     profile_path VARCHAR(255)
6 );
```

Listing 2: A unified People table

## 3 Connecting the Dots: Relationships

The database relational structure defines the connections between entities.

### 3.1 The Cast & Crew

Connecting `movies` to `people` required more than a simple link. We needed to know *what role* they played. For the **Cast**, we stored the `character_name` and an `order_index`. The index is crucial because it allows the frontend to validly display the "Top Billed" actors first, rather than a random alphabetical list. For the **Crew**, it's about the job and `department`. This granular detail allows us to answer interesting questions like "Which directors also acted in their own movies?"

```
1 CREATE TABLE IF NOT EXISTS cast_members (
2     credit_id VARCHAR(255) PRIMARY KEY,
3     movie_id INT,
4     person_id INT,
5     character_name VARCHAR(500),
6     order_index INT,
7     FOREIGN KEY (movie_id) REFERENCES movies(id),
8     FOREIGN KEY (person_id) REFERENCES people(id)
9 );
```

Listing 3: Linking People to Movies

## 3.2 Categorization (Genres, Keywords)

We handled categories like Genres and Keywords as Many-to-Many relationships. This means a movie can be both "Action" and "Comedy". We separated the names into their own lookup tables (`genres`, `keywords`) and used junction tables (`movie_genres`) to link them. This makes searching for "All Action movies" incredibly fast because it's just a simple index lookup.

# 4 User Interaction

## 4.1 Ratings

We built the `ratings` table. We decided to use a string-based `user_id` to align with our "Soft Session" strategy on the client. This table grows the fastest, so we kept it lightweight: just who, what, what score, and when.

```
1 CREATE TABLE IF NOT EXISTS ratings (
2     id INT AUTO_INCREMENT PRIMARY KEY,
3     movie_id INT,
4     user_id VARCHAR(255) NOT NULL,
5     rating FLOAT,
6     timestamp DATETIME DEFAULT CURRENT_TIMESTAMP,
7     FOREIGN KEY (movie_id) REFERENCES movies(id)
8 );
```

Listing 4: Capturing user votes

# 5 Binary Storage

## 5.1 Storing Images in the DB

We opted to store movie posters directly in the database using the `LONGBLOB` data type. While external storage (like S3) is a common pattern, storing images within the database ensures the dataset is self-contained. This simplifies backup and deployment procedures by keeping the entire application state in a single location.

```
1 CREATE TABLE IF NOT EXISTS movie_posters (
2     movie_id INT PRIMARY KEY,
3     image LONGBLOB,
4     FOREIGN KEY (movie_id) REFERENCES movies(id)
5 );
```

Listing 5: Self-contained image storage